

Econometrics Week7

William Schill

Final Assignment. Developed in Python and code hidden with jupyterutils. Code version can be submitted as well.

Each portion is separated into "PART X" and the reponses are in bold below the hidden or unhidden code portions.

```
In [3]: import numpy as np
import pandas as pd
from matplotlib.pyplot import *
import statsmodels.api as sma
import statsmodels.stats as sms
import statsmodels.tsa.stattools as smts
import seaborn as sns
import sklearn as skl
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
In [4]: from dfast.jupyterutils.togglecode import hideCode
hideCode()
```

Out[4]: Toggle Code.

Bring in the data and check it out. Then fully shift (lag) the data for future use and print to confirm.

```
In [5]: path = 'C:\\Users\\SchillW\\Documents\\Econ_Coursera\\Wk7\\'
df = pd.read_excel(path+'data.xls', index_col=0)
df.head(3)
```

Out[5]:

	GDP	GDPIMPR	LOGGDP	GrowthRate	li1	li2	T
Date							
1950Q1	94.300	NaN	4.546481	NaN	0	0	0
1950Q2	95.200	1.0	4.555980	0.009499	0	0	1
1950Q3	97.663	1.0	4.581523	0.025543	3	1	2

```
In [6]: ## FULLY SHIFT AND LAG THE DATA SET HERE
full = pd.concat([df, df.shift(1), df.shift(2)], axis=1, join='outer')
cols = []
for lg in range(0,3):
    for i in df.columns:
        if lg==0:
            cols.append(i)
        else:
            cols.append(i+'(-'+str(lg)+')')

full.columns = cols
```

Fully lagged dataframe of variables.

```
In [7]: print full.head(3)
```

	GDP	GDPIMPR	LOGGDP	GrowthRate	li1	li2	T	GDP(-1)	\
Date									
1950Q1	94.300	NaN	4.546481	NaN	0	0	0	NaN	
1950Q2	95.200	1.0	4.555980	0.009499	0	0	1	94.3	
1950Q3	97.663	1.0	4.581523	0.025543	3	1	2	95.2	

	GDPIMPR(-1)	LOGGDP(-1)	...	li1(-1)	li2(-1)	T(-1)	GDP(-2)	\
Date			...					
1950Q1	NaN	NaN	...	NaN	NaN	NaN	NaN	
1950Q2	NaN	4.546481	...	0.0	0.0	0.0	NaN	
1950Q3	1.0	4.555980	...	0.0	0.0	1.0	94.3	

	GDPIMPR(-2)	LOGGDP(-2)	GrowthRate(-2)	li1(-2)	li2(-2)	T(-2)
Date						
1950Q1	NaN	NaN	NaN	NaN	NaN	NaN
1950Q2	NaN	NaN	NaN	NaN	NaN	NaN
1950Q3	NaN	4.546481	NaN	0.0	0.0	0.0

[3 rows x 21 columns]

PART A - Loglikelihood Ratio Test

```
In [30]: try:
          del x, y
        except:
          pass

start = '1951Q1'
end = '2010Q4'

x = full[['li1(-1)', 'li2(-1)']]
xA = x[(x.index>=start)&(x.index<=end)]

y = full[['GDPIPR']]
yA = y[(y.index>=start)&(y.index<=end)]

modA = sma.Logit(endog=yA, exog=sma.add_constant(xA))
fitA = modA.fit()
print fitA.summary2()
```

Optimization terminated successfully.

Current function value: 0.559076

Iterations 5

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: 0.122
Dependent Variable:   GDPIPR                AIC:                274.3565
Date:                2017-03-08 12:05       BIC:                284.7984
No. Observations:    240                   Log-Likelihood:     -134.18
Df Model:            2                     LL-Null:           -152.76
Df Residuals:        237                   LLR p-value:        8.4833e-09
Converged:           1.0000                Scale:            1.0000
No. Iterations:      5.0000

-----
              Coef.    Std.Err.    z      P>|z|      [0.025    0.975]
-----
const      0.7288     0.1536    4.7454   0.0000     0.4278    1.0298
li1(-1)   -0.3719     0.0727   -5.1176   0.0000    -0.5143   -0.2294
li2(-1)   -0.1203     0.0377   -3.1936   0.0014    -0.1941   -0.0465
=====
```

Build the model from the data and match with the output on the test. Set up each specific model fit.

```
In [9]: fitA_Cli1 = sma.Logit(endog=yA, exog=sma.add_constant(xA[['li1(-1)']])).fit()
        fitA_Cli2 = sma.Logit(endog=yA, exog=sma.add_constant(xA[['li2(-1)']])).fit()
```

Optimization terminated successfully.

Current function value: 0.582277

Iterations 5

Optimization terminated successfully.

Current function value: 0.623002

Iterations 5

```
In [10]: loglikA = np.reshape(np.array([fitA.llnull, fitA_Cli1.llf, fitA_Cli2.llf,
fitA.llf]), (4,1))

loglik = np.reshape(np.array([-152.763, -139.747, -149.521, -134.178]), (4,1))
lldf = pd.DataFrame(loglik, index=['Const', 'Const+li1', 'Const+li2', 'All'], col
umns=['Loglikelihood'])
```

```
In [11]: print "The loglikelihood estimations are below. \n "
print "Loglikelihood:\n", loglikA
# print "Loglikelihood:\n", "\n", loglik, "\n \n LL Calc\n", loglikA
```

The loglikelihood estimations are below.

```
Loglikelihood:
[[-152.76340039]
 [-139.74657981]
 [-149.52046643]
 [-134.17823347]]
```

```
In [14]: from scipy.stats import chisqprob
def likelihood_ratio(llmin, llmax):
    return(2.0*(llmax-llmin))

LR = np.zeros((len(loglikA),len(loglikA)))
dof = LR.copy()
for i in range(4):
    for j in range(4):
        LR[i,j] = likelihood_ratio(loglikA[i],loglikA[j])
        if i==j:
            dof[i,j] = 0
        elif np.abs(j-i)<3:
            dof[i,j] = 1
        else:
            dof[i,j] = 2

# p = np.around(chisqprob(LR, dof), 3) # L2 has 1 DoF more than L1
p = np.around(chisqprob(LR, dof), 5)
pF = np.reshape(p[0,1:], (1,3))

print "P-Values for Loglikelihood Comparisons of Model \n"
pd.DataFrame(pF, columns = ['Const+li1', 'Const+li2', 'All Params'],
            index = ['vs Null'])
```

P-Values for Loglikelihood Comparisons of Model

```
Out[14]:
```

	Const+li1	Const+li2	All Params
vs Null	0.0	0.01087	0.0

(a):

The matrix above reflects that the model is valid as all of the p-values are significant to the 1% level with the Constant and li2 model being significant to the 5% level.

Additionally, this can be confirmed by examining the llr_pvalue features of the statsmodels models as seen below:

```
In [16]: lrsma = np.array([fitA_Cli1.llr_pvalue, fitA_Cli2.llr_pvalue,
                        fitA.llr_pvalue], ndmin=2)
lrsma = pd.DataFrame(lrsma.T, columns=['Likelihood Ratio P-Values'])
lrsma.index = ['Const+li1', 'Const+li2', 'Full Model']
np.around(lrsma, 7)
```

```
Out[16]:
```

	Likelihood Ratio P-Values
Const+li1	3.000000e-07
Const+li2	1.087350e-02
Full Model	0.000000e+00

PART B

The calculation for Mcfadden's is : $Mcfaddens R^2 = 1 - LL_Model / LL_Null$.

```
In [17]: x = full[['li1(-1)','li2(-1)','li1(-2)','li2(-2)']]
xB = x[(x.index>=start)&(x.index<=end)]

y = full[['GDPIMPR']]
yB = y[(y.index>=start)&(y.index<=end)]

modB = sma.Logit(endog=yB, exog=sma.add_constant(xB))
fitB = modB.fit()

fitB_C111121 = sma.Logit(endog=yB, exog=sma.add_constant(xB[['li1(-1)','li2(-1)']])).fit()
fitB_C111122 = sma.Logit(endog=yB, exog=sma.add_constant(xB[['li1(-1)','li2(-2)']])).fit()
fitB_C112121 = sma.Logit(endog=yB, exog=sma.add_constant(xB[['li1(-2)','li2(-1)']])).fit()
fitB_C112122 = sma.Logit(endog=yB, exog=sma.add_constant(xB[['li1(-2)','li2(-2)']])).fit()

## Doing it this way left us with different results than expected.
```

```
Optimization terminated successfully.
      Current function value: 0.540830
      Iterations 6
Optimization terminated successfully.
      Current function value: 0.559076
      Iterations 5
Optimization terminated successfully.
      Current function value: 0.558860
      Iterations 5
Optimization terminated successfully.
      Current function value: 0.543106
      Iterations 6
Optimization terminated successfully.
      Current function value: 0.543588
      Iterations 6
```

```

In [18]: LL_null = -152.763
         LL_cl11121 = -134.178
         LL_cl11122 = -134.126
         LL_cl12121 = -130.346
         LL_cl12122 = -130.461

         LLB = pd.DataFrame(np.reshape(np.array([-152.763, -134.178, -134.126,
         -130.346, -130.461]),(1,5)))
         LLB.columns = ['LL_null','LL_cl11121','LL_cl11122','LL_cl12121','LL_cl12122']

         MR2 = pd.DataFrame(np.reshape(np.array([1.0 -
         LLB['LL_cl11121']/LLB['LL_null'],
                                     1.0 -
         LLB['LL_cl11122']/LLB['LL_null'],
                                     1.0 - LLB['LL_cl12121']/LLB['LL_null'],
                                     1.0 -
         LLB['LL_cl12122']/LLB['LL_null']]),(1,4)))
         MR2.columns = ['LL_cl11121','LL_cl11122','LL_cl12121','LL_cl12122']

         print "Log Likelihoods : \n", LLB
         print "\n McFaddens R Squared : \n", MR2

```

Log Likelihoods :

	LL_null	LL_cl11121	LL_cl11122	LL_cl12121	LL_cl12122
0	-152.763	-134.178	-134.126	-130.346	-130.461

McFaddens R Squared :

	LL_cl11121	LL_cl11122	LL_cl12121	LL_cl12122
0	0.121659	0.121999	0.146744	0.145991

(b):

Based on the calculated McFaddens R^2 above, the most optimal model would be to go with the $Constant + li1(-2) + li2(-1)$ model. But not by very much.

PART C

Probability of economic growth over evaluation sample.

```
In [31]: ## The model was created above:
## It should be noted that they do not accurately specify the time constraints
for these models!!!!
print fitB_Cl12l21.summary2()
```

Results: Logit

```
=====
Model:                Logit                Pseudo R-squared: 0.147
Dependent Variable:   GDPIMPR                AIC:                266.6909
Date:                2017-03-08 12:05        BIC:                277.1328
No. Observations:    240                    Log-Likelihood:    -130.35
Df Model:            2                      LL-Null:          -152.76
Df Residuals:        237                    LLR p-value:       1.8366e-10
Converged:           1.0000                  Scale:           1.0000
No. Iterations:      6.0000

-----
              Coef.      Std.Err.      z      P>|z|      [0.025      0.975]
-----
const        0.7457      0.1573      4.7397    0.0000     0.4373     1.0540
li1(-2)     -0.4287      0.0763     -5.6175    0.0000    -0.5783    -0.2791
li2(-1)     -0.1312      0.0386     -3.3994    0.0007    -0.2068    -0.0556
=====
```



```

In [20]: x = full[['li1(-2)','li2(-1)']]
        xC = x[x.index>end]

        predC = fitB_C112121.predict(sma.add_constant(xC))
        predCTable = fitB_C112121.pred_table(threshold=0.5)
        rtot = np.sum(predCTable,axis=1)
        tot = np.sum(rtot)

        predCPRT = pd.DataFrame(predCTable/tot)
        predCPRT.columns = ['yhat=0','yhat=1']
        predCPRT.index = ['y=0','y=1']

        predSum = pd.DataFrame(rtot)
        predSum.columns=['Sum']
        predSum.index = ['y=0','y=1']
        prt = pd.concat([predCPRT, predSum], axis=1, join='outer')

        print "Prediction Realization Table: \n"
        print prt, "\n"
        print "Sums :"
        print np.sum(prt,axis=0), "\n"
        print "Hit Rate :"
        print prt.iloc[0,0] + prt.iloc[1,1]

```

Prediction Realization Table:

	yhat=0	yhat=1	Sum
y=0	0.133333	0.2000	80.0
y=1	0.104167	0.5625	160.0

Sums :

yhat=0	0.2375
yhat=1	0.7625
Sum	240.0000

dtype: float64

Hit Rate :

0.695833333333

(c): THE PREDICTION REALIZATION TABLE AND HIT RATE ARE PRINTED ABOVE.

The hit rate is 69.58 and the probability of economic growth is approximately 76.25 in total.

The hit rate represents the accuracy of the model. With 1 being that the model is increasing we are likely to continue to see economic growth.

PART D

ADF Test for Log GDP. The ADF test is formulated and in Python you can print out the model results as seen below.

```
In [32]: x = full[['LOGGDP']]
xD = x[(x.index>=start)&(x.index<=end)]

loggdp_ADF = sma.tsa.stattools.adfuller(xD['LOGGDP'], maxlag=1,
                                         autolag=None, regression='ct',
                                         regresults=True) #c for constant only,
                                         no trend
print 'ADF Test \n'
print 'Statistic :', loggdp_ADF[0], ', P-value :', loggdp_ADF[1], '\n'
print 'Confidence Levels :', loggdp_ADF[2], '\n'
print 'Model \n'
print loggdp_ADF[3].resols.summary2(xname=['LOG GDP lag1', 'Diff LOG GDP
lag1', 'Constant', 'Trend'])
### CONFIRMED THE ORDER OF CONST TREND
```

ADF Test

Statistic : -2.51821089277 , P-value : 0.318918118592

Confidence Levels : {'5%': -3.4290999471622556, '1%': -3.9973200578432064, '10%': -3.1379848180498104}

Model

Results: Ordinary least squares

```
=====
Model: OLS Adj. R-squared: 0.393
Dependent Variable: y AIC: -1888.3395
Date: 2017-03-08 12:05 BIC: -1874.4504
No. Observations: 238 Log-Likelihood: 948.17
Df Model: 3 F-statistic: 52.20
Df Residuals: 234 Prob (F-statistic): 7.23e-26
R-squared: 0.401 Scale: 2.0630e-05
```

```
-----
              Coef. Std.Err. t P>|t| [0.025 0.975]
-----
LOG GDP lag1 -0.0204 0.0081 -2.5182 0.0125 -0.0364 -0.0044
Diff LOG GDP lag1 0.6325 0.0509 12.4338 0.0000 0.5323 0.7328
Constant 0.0956 0.0375 2.5512 0.0114 0.0218 0.1695
Trend 0.0001 0.0000 2.4979 0.0132 0.0000 0.0001
-----
```

```
-----
Omnibus: 23.009 Durbin-Watson: 2.012
Prob(Omnibus): 0.000 Jarque-Bera (JB): 36.479
Skew: 0.578 Prob(JB): 0.000
Kurtosis: 4.530 Condition No.: 23965
=====
```

* The condition number is large (2e+04). This might indicate strong multicollinearity or other numerical problems.

(d):

The coefficient value for LOG GDP at 1 Lag is -0.0204 and the subsequent p-value is 0.0125 which is significant to the 5% level. I do not think this is a necessary statistic to point out but it was in previous tests.

Based on the output listed above using the statsmodels ADF test, the test statistic is -2.518 and its corresponding p-value is 0.3189. It is very far right of even the 10% confidence level and would thus imply that we cannot reject H_0 and the variable is NOT stationary.

PART E

Modeling the Growth Rate at 1 Lag with varying lags of the First and Second Leading Indicators. Only the R Squares from the models are presented.

```

In [23]: x1 = full[['GrowthRate(-1)', 'li1(-1)', 'li2(-1)']]
x11 = x1[(x1.index>=start)&(x1.index<=end)]

x2 = full[['GrowthRate(-1)', 'li1(-2)', 'li2(-1)']]
x12 = x2[(x2.index>=start)&(x2.index<=end)]

x3 = full[['GrowthRate(-1)', 'li1(-1)', 'li2(-2)']]
x21 = x3[(x3.index>=start)&(x3.index<=end)]

x4 = full[['GrowthRate(-1)', 'li1(-2)', 'li2(-2)']]
x22 = x4[(x4.index>=start)&(x4.index<=end)]

y1 = full[['GrowthRate']]
y11 = y1[(y1.index>=start)&(y1.index<=end)]
y12 = y11.copy()
y21 = y11.copy()
y22 = y11.copy()

modE11 = sma.OLS(endog=y11, exog=sma.add_constant(x11)).fit()
modE12 = sma.OLS(endog=y12, exog=sma.add_constant(x12)).fit()
modE21 = sma.OLS(endog=y21, exog=sma.add_constant(x21)).fit()
modE22 = sma.OLS(endog=y22, exog=sma.add_constant(x22)).fit()

r2 = pd.DataFrame([modE11.rsquared, modE12.rsquared, modE21.rsquared, modE22.rsquared])
r2.index = ['mod11', 'mod12', 'mod21', 'mod22']
r2.columns = ['R Squares']
r2

```

Out[23]:

	R Squares
mod11	0.507975
mod12	0.477193
mod21	0.507665
mod22	0.477130

```

In [24]: modE11Coefs = pd.DataFrame(modE11.params,
                                     columns=['Coefficients of Model k1=k2=1'])
modE11Coefs

```

Out[24]:

	Coefficients of Model k1=k2=1
const	0.001737
GrowthRate(-1)	0.461579
li1(-1)	-0.001023
li2(-1)	-0.000149

(e):

The results above show the best R^2 is indeed 0.5079 for the model where $k_1 = k_2 = 1$ (Mod11) barely beating out Mod21. The coefficients are then listed below that where the Growth Rate lagged variable has a coefficient of 0.461579. Such a large value in comparison to the other coefficients might indicate that we have a misspecification in our model.

PART F

Breusch Godfrey Test to examine the lag structure and see if the model was misspecified.

```
In [25]: bgall = sms.diagnostic.acorr_breush_godfrey(modE11, nlags=1, store=True)
bg = bgall[0:4]
modE11_BG = pd.DataFrame(np.reshape(np.array(bg), (1,4)),
                          columns=['LM Stat', 'LM p-val', 'Fvalue', 'Fval:p-val'],
                          index=['BG Test for Model k1=k2=1'])
modE11_BG
```

```
Out[25]:
```

	LM Stat	LM p-val	Fvalue	Fval:p-val
BG Test for Model k1=k2=1	0.230366	0.631253	0.225783	0.63511

(f):

The Breusch Godfrey test for the model where $k_1 = k_2 = 1$ at 1 Lag shows a high p-value for the test statistic of 0.230366. This test has H_0 that there is no serial correlation and with a low p-value we reject the Null. Thus with a p-value of 0.631253 we cannot reject the null and there is no serial correlation.

PART G

Forecast for the Growth Rate.

```
In [26]: x = full[['GrowthRate(-1)', 'li1(-1)', 'li2(-1)']]
xG = x[x.index>end]
xGR = pd.DataFrame(full[full.index>end]['GrowthRate'])

predG = modE11.predict(sma.add_constant(xG))

predGdf = pd.DataFrame(predG, index=xG.index)
outDf = pd.concat([predGdf, xGR], axis=1, join='outer')
outDf.columns = ['Prediction', 'Actual']
```

```
In [28]: fig2, ax=subplots(1,1,figsize=(12,6))
outDf.plot(ax=ax, title='Growth Rate Examination')
```

```
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0xb599710>
```



```
In [29]: maeG = np.sum(np.absolute(predGdf.values-xGR.values))
rmseG = np.sqrt(((predGdf.values-xGR.values) ** 2).mean())
print "FIT G : RMSE \ MAE = ", rmseG, " \ ", maeG, "\n"
```

```
FIT G : RMSE \ MAE = 0.00315576342885 \ 0.0528065895385
```

(g):

The above is the time series representation of the prediction and actuals over the 'evaluation' period. The RMSE is 0.003156 and the MAE is 0.0528.

The prediction is significantly smoother than the actuals which is to be expected and based on appearance and error the prediction model fits the actuals very well.

```
In [72]: from dfast.jupyterutils.togglecode import hideCode
hideCode()
```

```
Out[72]: Toggle Code.
```

```
In [ ]:
```